

CSE 486/586 Distributed Systems Consistency --- 2

Steve Ko
Computer Sciences and Engineering
University at Buffalo

CSE 486/586, Spring 2013

Linearizability vs. Sequential Consistency

- Both care about giving an **illusion of a single copy**.
 - From the outside observer, the system should (almost) behave as if there's only a single copy.
- Linearizability cares about **time**.
- Sequential consistency cares about **program order**.
- We need to look deeper into both concepts to understand the difference.

CSE 486/586, Spring 2013

2

Linearizability

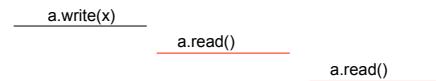
- Linearizability
 - Should provide the behavior of a single copy
 - A read operation returns the most recent write, regardless of the clients.
 - All subsequent read ops should return the same result until the next write, regardless of the clients.
- "The most recent" & "all subsequent"
 - Determined by time.
- Complication
 - In the presence of concurrency, read/write operations overlap.

CSE 486/586, Spring 2013


3

Linearizability Complications

- Non-overlapping ops: time-based clear-cut ordering



- Overlapping ops: not clear-cut with time



CSE 486/586, Spring 2013

4

Linearizability Complications

- Non-overlapping ops: time-based clear-cut ordering
 - Global time determines "most recent write" & "subsequent reads"
- Overlapping ops: not clear-cut with time
 - The system needs to provide an ordering of ops.
 - The ordering should give an illusion that it has a single copy.
- I.e., some ordering of operations where:
 - A read returns the result of the most recent write.
 - Once the result of the write becomes visible, all subsequent reads return the same result until the next write becomes visible.

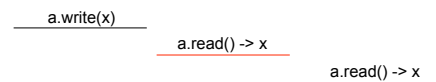
CSE 486/586, Spring 2013

5

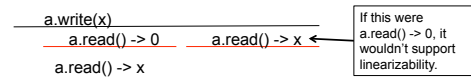
Linearizability Examples



- Example 1



- Example 2



CSE 486/586, Spring 2013

6

Linearizability Examples

- Example 3

```

a.write(x)
a.read() -> x    a.read() -> x
a.read() -> y
a.write(y)
    
```

CSE 486/586, Spring 2013

7

CSE 486/586 Administrivia


- PA3 deadline: 3/29 (Friday)
- Anonymous feedback form still available.
- Please come talk to me!

CSE 486/586, Spring 2013

8

Chain Replication


- One technique to provide linearizability



CSE 486/586, Spring 2013

9

Passive (Primary-Backup) Replication



- **Request Communication:** the request is issued to the primary RM and carries a unique request id.
- **Coordination:** Primary takes requests atomically, in order, checks id (resends response if not new id.)
- **Execution:** Primary executes & stores the response
- **Agreement:** If update, primary sends updated state/result, req-id and response to all backup RMs (1-phase commit enough).
- **Response:** primary sends result to the front end

CSE 486/586, Spring 2013

10

Linearizability vs. Sequential Consistency

- Both care about giving **an illusion of a single copy**.
 - From the outside observer, the system should (almost) behave as if there's only a single copy.
- Linearizability cares about **time**.
- Sequential consistency cares about **program order**.

CSE 486/586, Spring 2013

11

Sequential Consistency

- Sequential consistency
 - Should provide the behavior of a single copy
 - A read operation returns the most recent write, regardless of the clients.
 - All subsequent read ops should return the same result until the next write, regardless of the clients.
- "The most recent" & "all subsequent"
 - Ops within the same client: Determined by time (program order)
 - Ops across clients: **Not** determined by time, i.e., we can re-order them.
 - I.e., we just need to preserve the program order

CSE 486/586, Spring 2013

12

Sequential Consistency

- To the outside observer, the system needs to provide a global ordering of operations where:
 - It works like a single copy.
 - The ordering of ops coming from the same client is preserved.
- Linearizability vs. sequential consistency
 - With sequential consistency, the system has freedom as to how to interleave operations coming from **different clients**, as long as the ordering from each client is preserved.
 - With linearizability, the interleaving across all clients is pretty much determined already based on time.

CSE 486/586, Spring 2013

13


Sequential Consistency Examples

- Example 1
 - P1: a.write(A)
 - P2: a.write(B)
 - P3: a.read()->B a.read()->A
 - P4: a.read()->B a.read()->A
- Example 2
 - P1: a.write(A)
 - P2: a.write(B)
 - P3: a.read()->B a.read()->A
 - P4: a.read()->A a.read()->B

CSE 486/586, Spring 2013

14

Active Replication



- Request Communication:** The request contains a unique identifier and is multicast to all by a reliable totally-ordered multicast.
- Coordination:** Group communication ensures that requests are delivered to each RM in the same order (but may be at different physical times!).
- Execution:** Each replica executes the request. (Correct replicas return same result since they are running the same program, i.e., they are replicated protocols or replicated state machines)
- Agreement:** No agreement phase is needed, because of multicast delivery semantics of requests
- Response:** Each replica sends response directly to FE

CSE 486/586, Spring 2013

15

Summary

- Linearizability
 - The ordering of operations is determined by time.
 - Primary-backup can provide linearizability.
 - Chain replication can also provide linearizability.
- Sequential consistency
 - The ordering of operations preserves the program order of each client.
 - Active replication can provide sequential consistency.

CSE 486/586, Spring 2013

16

Acknowledgements

- These slides contain material developed and copyrighted by Indranil Gupta (UIUC).

CSE 486/586, Spring 2013

17